ATLAS Mesh Design

Furniture that has the shape of the furniture

4 October 2026. Branch `main`. For the team lead and the three engineers; Appendix A is for the agents. A design to build now. The decisions in section 2 were taken by the team lead on 4 October 2026. Facts about the code were read from this tree; anything that could not be checked says so. The same text is in `ATLAS_Mesh_Design.md`.

Corrected on 4 October 2026 after the first build, so that it says what the code does. Nothing in it has run on a phone, and the real generator (TripoSG) has never run. When this text was corrected, the Forge's fit, attempts, colour per face and light copy existed as separate modules (`forge/profile.py`, `forge/colour.py`) and were not yet wired into clean-up (section 5.3).

1 The problem and the outcome

What is seen today. Scanned objects stand in the right place and have the right size. Their shape is wrong: a chair is a generic chair that often faces the wrong way, some objects are plain boxes, and only some look right.

Why. Position and size come from RoomPlan's box, which is good. Shape comes from three places that all fall short:

Seen	Cause	Where
Every chair is the same blocky chair	The real generator, TripoSG, has never run. What is shown is the built-in set: box composites of 22 to about 100 triangles, one per category	`station/forge/forge/generators/triposg.py`, `tools/gen_fallback_set.py`, `ATLAS_Whats_Left.md` item 9
It faces the wrong way	An object's yaw is the forward axis of RoomPlan's box, and nothing in our code decides which side of the box is the front. The built-in chair has its back on local −Z whatever the real chair does	`Scan/Routes/ObjectBoxMerge.cs:54`
A generated mesh would face the wrong way too	The Forge turns a mesh by comparing its bounds with the box. That tells 0° from 90° only when the footprint is not square, and never 0° from 180°. The phone already sends each photo's camera pose (`camToBox`, `azimuthDeg`); the Forge does not read them	`forge/cleanup.py:236`, `forge/models.py:35`
A plain box	No template for `unknown` and `other`; only the 20 largest objects ever leave their box; a template that splits past 1,500 triangles is dropped	`tools/gen_fallback_set.py`, `forge/budget.py`, `FallbackMeshBuilder.cs:41`
One flat colour	Clean-up picks one palette cell per object from the photo's dominant colour	`forge/cleanup.py:502`
No doors or windows	The RoomPlan wall bridge is written and switched off	`Scan/Routes/WallBridge.cs:53`

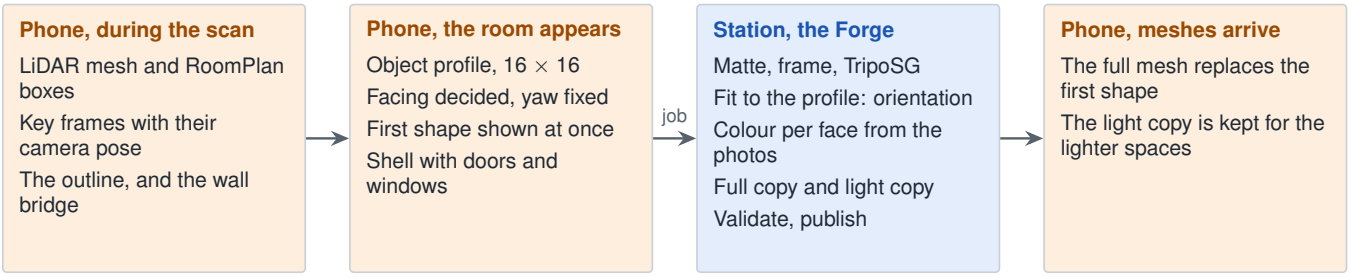
The outcome. When this design is built:

1. Every scanned object has a shape of its own kind the moment the room appears, facing the way the real one faces.
2. Within about a minute per object, that shape is replaced by a mesh generated from the photo of that very object.
3. The mesh sits where the real object is, to the accuracy the scan already has. Accuracy comes first: nothing in this design moves an object or changes its box.
4. Objects show the colours of their parts, in the palette.
5. A room may hold up to 32 generated meshes and 120,000 furniture triangles, with a lighter copy for the spaces that cannot afford that.
6. Walls have their real height, doors and windows.

2 Decisions

#	Subject	Decision	Was
M1	Purpose	A design to build now, centred on TripoSG producing real furniture shapes	–
M2	Priority	Accuracy first; looks are improved only where accuracy is not put at risk	–
M3	Shape	Two layers: a first shape at once on the phone, then the TripoSG mesh	Box, then mesh or built-in
M4	First shape	Named objects: the category's shape fitted to the scan. Unnamed objects: a shape cut from the scan	Fixed template, or a box
M5	Facing	Decided from the scan on the phone; the Forge also uses the scan and the photo's camera pose	RoomPlan's box axis
M6	Disagreement	The generated mesh is always shown, whatever the scan says. The scan chooses its orientation; it never rejects it	–
M7	Attempts	One generation per object from the best photo. Only when the fit to the scan is poor, the other photos are tried, and the best fit is shown	One, and one retry on failure
M8	Colour	Per-face palette colours projected from the scan photos. No texture, the one shader	One palette cell per object
M9	Triangles	Room total 120,000; per object 3,000 to 15,000; up to 32 objects	60,000; 3,000 to 5,000; 20
M10	Levels	The Forge returns a full and a light copy. Home and Curved Walk show the full one; the other spaces show the light one until measured	One copy
M11	Room shell	The outline-based shell stays. The RoomPlan wall bridge is turned on for wall heights, doors and windows	Bridge off
M12	Classification	Stays Apple's (ARKit surface classes, RoomPlan categories). No classifier of our own	Same
M13	Ray casting	Used inside the mesh pipeline (section 6.1). A ray-cast renderer is a future option, sketched in section 6.2	Not used
M14	Proof	Numbers from a recorded team scan, plus a walk-through on the phone. The record stays on the Station and never enters git	–
M15	Unchanged	The hosted third-party fallback stays off. The Forge queue stays largest object first	Same

3 How it works after this design



Three things are new and carry most of the design:

- **The object profile** (section 5.2): a 16 × 16 grid of heights of the scanned surface inside the object's box. It is made by casting rays straight down onto the LiDAR mesh. It is tiny, it travels in the RoomModel and in the Forge job, and it is what decides facing, fits the first shape and orients the generated mesh.
- **The fit** (section 5.3): the Forge makes the same profile from the generated mesh and compares the two. The orientation with the smallest difference wins.
- **Two levels** (section 5.4): a full and a light copy of every mesh, so the budget can rise without putting any space over its limit.

4 Work areas

For people. Each area is one piece of work with one owner lane. The tasks behind each area are in Appendix A.

Area	Delivers	Lane	Waits for	Proved by
A · TripoSG runs	A real mesh from a real photo, on the Station, in the Forge	D	–	M2
B · Facing and first shapes	The object profile; the facing rule; fitted and scan-cut first shapes; nothing is a plain box	C, D	–	M1
C · Fit, colour and levels	The Forge orients by the scan, tries again when the fit is poor, colours per face, and returns two copies	D	A (real meshes), B (the profile), D (the targets)	M2
D · The furniture budget	The new triangle rule everywhere it is written; per-space level; the phone measurement	A, C, D	–	M3
E · Doors and windows	The wall bridge built into the app and switched on	C	–	M4
F · Proof	The recorded scan, the numbers script, the phone checklist	D, lead	A to E	M1 to M4

Areas A, B, D and E do not wait for each other and can be built at the same time. Area C starts on the procedural generator and the fixtures, and is finished on real meshes once A is through. Area F closes the work.

5 The design

5.1 Area A · TripoSG runs on the Station

The adapter (`generators/triposg.py`) was written from the public repository without ever being run. Its calls were checked against that repository on 4 October 2026 (last upstream commit 18 April 2025; the links are in Appendix C). Nothing was run. “Read” means read in the source; “inferred” means worked out from it.

Subject	What the public repository says	Basis	What the Forge does
The calls	The import path, `from_pretrained(dir).to("cuda", float16)`, the call and `.samples[0]` as `(vertices, faces)` are as the adapter has them	Read	Nothing to change
Framing the picture	The picture is cut to the matte’s bounding box, put on white and padded to a square by 10 % of its long side on every side. The model’s own processor then resizes the short edge to 256 and takes the central 224 × 224, so a picture that is not square loses its sides	Read	The same framing (below). Today the adapter only puts the crop on white
Matting	Always matted. The model was trained on a foreground that fills 90 % of the picture	Read	Matting becomes required. We use `rembg`; upstream’s RMBG-1.4 has a licence of its own
Camera angles	Trained on views from the front half of the object, from 15° below to 30° above	Read (paper)	Photo selection prefers such views (below)
The decoder on Windows	The default decoder uses `diso`. One open report: on Windows 10 with Python 3.11 it fails with a CUDA error and the pipeline then crashes on an empty result	Read; the way round it is inferred and untested	On that failure the adapter tries again with `use_flash_decoder=False`. `availability()` also checks that `diso` imports
Installing	TripoSG’s `requirements.txt` pins numpy 1.22.3, which has no build for Python 3.11, the Forge’s version. `transformers` is not pinned: 4.57.0 fails to load the model, 4.56.2 works, the official demo pins 4.49.0	Read	The Forge keeps its own pinned list, `requirements-triposg.txt`. The README’s install line, which fails as written, is corrected
The weights	7.95 GB. A folder that was only partly downloaded breaks loading	Read	The adapter looks for `model_index.json`, not for the folder

Subject	What the public repository says	Basis	What the Forge does
The result	Vertices inside a cube of ± 1.005, with no real-world scale. Shape only: no colour, no texture. The authors' texturing route needs more memory than this card has and does not run on Windows	Read	Scale to the box as today. Colour by our own projection (section 5.3)
The result's frame	Not documented. The code around it suggests +Y up and the front towards +Z, in the object's own frame and not the camera's	Inferred	Measured on the first run (below)
The result's size	Surfaces are extracted on a 505^3 grid; the official demo cuts to 100,000 faces. Raw counts are not published and are probably several hundred thousand	Read; the count is inferred	Clean-up decimates first, as today, and its time-out rises from 20 s to 60 s
Memory and time	"At least 8 GB". One user reports about 9 GB used on a 12 GB card, another about 30 s per mesh. No official figure, and no report on an RTX 50-series card	Read; user reports	Check M2 measures both on the Station
Licences	Code and weights are MIT. `diso` is non-commercial. The decoder taken from FlashVDM is under Tencent's community licence	Read	Fine for the demo; to be looked at again before any commercial use

Set-up. PyTorch with the cu128 wheels (builds for Python 3.11 exist from 2.7.0), the TripoSG checkout, `requirements-triposg.txt`, `requirements-gpu.txt`, `FORGE_TRIPOSG_REPO`. `diso` is built from source, so the Station needs the Microsoft build tools and the CUDA toolkit with `CUDA_HOME` set; without `CUDA_HOME` it builds without its CUDA part and reports no error. The installed versions and the TripoSG commit go into `docs/versions.md`.

What the Forge's first round established beyond the research above (`station/forge/README.md`, "First run on the Station"; nothing here was run on the Station):

- **The lock file did not install on Windows.** `requirements.lock` pinned `uvloop` with no platform marker, and `uvloop` has no Windows build, so `pip install -r requirements.lock` would have failed on the Station. It is now compiled with `uv pip compile --universal`: `uvloop` is skipped on Windows, `colorama` is added there, and no version changed.
- **`diso` cannot be in a requirements file.** Its build script imports PyTorch and pip builds in an isolated environment that has none (seen: `ModuleNotFoundError: No module named 'torch'`). It is installed by its own command, after PyTorch: `pip install setuptools wheel`, then `pip install --no-build-isolation diso==0.1.4` with `CUDA_HOME` set. `opencv-python` and `pymeshlab`, which only TripoSG's scripts use, are left out of the list.
- **The pinned versions.** `torch==2.7.1` and `torchvision==0.22.1` from the cu128 index, `diffusers==0.33.1`, `transformers==4.49.0`, `huggingface-hub==0.30.2`, `scikit-image==0.25.2`, `scipy==1.15.2`, `peft==0.15.2`, `diso==0.1.4`, `rembg[gpu]==2.0.69`, TripoSG commit `fc5c40990181e2a756c4e0b1c2f4d6b5202faf8c`. The lists were resolved with `uv pip compile` for Python 3.11 on Windows, and every pin has a Windows wheel except `antlr4-python3-runtime`, which is pure Python.
- **What was tried, and where.** In a throw-away Python 3.11 environment on the build Mac, with PyTorch for the processor, the TripoSG package imported with `diso` replaced by a stub, and `rembg` cut four of TripoSG's example pictures, which the Forge framed. No weights were downloaded and nothing was generated. The Windows commands of the run-book are written from reading, not from running.
- **What exists to run it.** The adapter's fallback decoder, the `model_index.json` and `diso` checks, the model loaded when the service starts, the run-book, `tools/check_station.py` (checks the install) and `tools/first_run.py` (posts a room and one photo, follows the job, validates the GLB and reads the frame).

The input picture. The Forge prepares the crop the way the repository's own script does:

1. Matte the crop. Matting becomes required for the `triposg` generator: `GET /health` reports `generatorReady: false` when no matting model is installed, instead of sending unmatted room photos to the model.
2. Cut the picture to the matte's bounding box, put it on white, and pad it to a square by 10 % of its long side on every side.
3. Hand that picture to the pipeline. The original crop and its matte are kept in memory for the colour step (section 5.3).

The views. Facing is known when the photos are chosen (section 5.2), so selection can prefer what the model was

trained on: a view's score is halved when the camera stands in the back half of the object, or more than 30° above its centre. A view from behind is still taken when there is no other.

The output frame. Which way is up, and where the mesh's front is, is measured once on the Station with a test photo of an object that is clearly not symmetric (a chair seen from the front-left). The answer is set as `triposg_up` and `triposg_front` in `forge/config.py`; they are settings, read from `FORGE_TRIPOSG_UP` and `FORGE_TRIPOSG_FRONT` (defaults `Y` and `unknown`) and checked at start-up, and nothing reads them yet because the fit is not wired into clean-up. `tools/first_run.py` reads the frame from a finished job. `triposg_front` is `canonical` (the front is on the mesh's own +Z, which after section 5.2 is also the box's front), `view` (the front follows the camera) or `unknown`. Section 5.3 uses the scan to choose the orientation in any case; this setting only breaks ties.

Loading. The model is loaded when the service starts, not on the first job, so the first object does not pay for it. `GET /health` reports `generatorReady` only when the load succeeded.

Limits. Generation time-out stays 90 s. The plan's proposal of 60 s per object or less (`P1-D08`) is the figure check M2 measures.

5.2 Area B · The object profile, facing and first shapes

The object profile

For each object box of a scan that has a real LiDAR mesh (`ScanSnapshot.MeshFromPlanes` is false):

- **Grid.** 16 × 16 cells over the box's footprint in the box's local frame. Cell (i, j) covers local x from $-w/2 + i\,w/16$ and local z from $-d/2 + j\,d/16$. Index $= 16j + i$, row 0 at the lowest z, as the room's distance field.
- **Rays.** Four rays per cell (a 2 × 2 pattern at the quarter points), 1,024 per object. Each starts 0.05 m above the top of the box and points straight down. It is tested against the scan triangles that overlap the box, floor-class triangles left out. A hit counts when it lies from 0.02 m above the base to the ray's start.
- **Value.** The highest hit of the cell's four rays, as a byte: `q = clamp(floor(255 * height / h + 0.5), 1, 255)`. `0` means no ray of the cell hit anything.
- **Coverage.** Cells that are not 0, divided by 256.
- **Where it lives.** `objects[k].profile` in the RoomModel: the 256 bytes as base64 (344 characters). The key is optional. It is absent for route 1, for a phone without LiDAR, for fixtures, and for an object in which no ray hit anything (coverage 0); a profile is never stored as 256 zeros. It is not part of the geometry hash. 64 objects add about 23 kB to a file whose limit is 300 kB.
- **Where it is made.** `Atlas.Scan`, engine-free, inside `ScanSolver.Solve` on its worker thread, which already holds the mesh, its classes and the boxes.

Facing

Facing is decided once, when the RoomModel is built, by the first rule in the object's list that gives an answer. After it, **local +Z is the front of every object**: the yaw is turned by whole quarter turns, `size.w` and `size.d` are swapped on an odd turn, and the profile is turned with them. No reader of the RoomModel needs to change.

Rule	Answer	Accepted when
`profile`	The back is the side whose edge strip (the outer 3 of 16 rows) is highest on average; the front is opposite	That strip has at least a quarter of its cells seen, and its mean height exceeds the next side's by 0.15 of the box height
`table`	The front is the side whose outward normal points most towards the centre of the nearest table	A table's box lies within 1.0 m of the object's box
`clearance`	The front is the side with the most free floor in front of it: the room's distance field sampled 0.30 m outside the middle of each side	The best side exceeds the opposite side by 0.10 m
`box`	The box's own axis, as today	Always

Categories	Rules, in order
chair	`profile`, `table`, `clearance`, `box`
couch, bed, toilet	`profile`, `clearance`, `box`
storage, screen, refrigerator, oven, dishwasher, washerDryer, stove, fireplace, sink, bathtub	`clearance`, `box`
table, lamp, plant, stairs, keyboard, mouse, laptop, other, unknown	`box`

Ties in `clearance`. An object's own box limits what any open side can read: the probe is 0.30 m out but the free region stops one cell (0.05 m) short, so every open side reads at most 0.25 m. A deep appliance against a wall then reads the same on all three open sides, and "the most free floor" is an exact three-way tie, which taken literally goes to the lowest side: right only when RoomPlan's axis happens to point into the room, and `box` otherwise, which is the case the rule exists for. So the sides that read within one cell (0.05 m, `ClearanceTie`) of the best count as equally free, and of those the side that exceeds its opposite side by most is taken. The 0.10 m acceptance is unchanged, and an exact tie that remains (a square box in a corner) goes to the lowest quarter turn.

Two readings that the table leaves open. In `profile`, "the next side's" mean is the highest mean of the other three strips (a tie gives no answer); the strips share their corner cells, so a scan that saw only the back gives no answer; the quarter-seen test applies to the highest strip only, so a strip that is highest on a few cells blocks the rule and the next one answers. In `table`, "within 1.0 m" is the distance between the two footprints on the floor, 0 when they overlap; of two tables equally near the nearer centre wins, and a table that is the object itself gives no answer. In both rules "exceeds by" means at least.

The rule that answered is written to `objects[k].facing` (optional key; absent means `box`). Check M1 counts it.

Facing can be switched off by a scan option, `ScanOptions.DecideFacing` (true in `ScanOptions.Default`). With it off the profile is still computed, because the first shapes and the Forge read it, and only the turning and `facing` are left out. In this tree `ScanFlow` does not yet hand the option to `ScanSolver.Solve`, so in the app facing is always on; only a direct call of `ScanSolver.Solve` or `RoomModelBuilder.Build` can switch it off.

First shapes

The asset client shows a first shape for every object as soon as the room content is built, not only when the Forge is away. The mesh state stays `builtIn`; the built-in limit of 1,500 triangles and the 0.25 m edge rule stay.

Object	First shape
chair, couch, bed	The category's template with its measures taken from the profile: seat or mattress height is the median height of the seen cells in the front half, the back stands on local $-Z$ at the height of the box. Defaults are the template's own when the cells are unseen. The height is kept within 0.25 to 0.70 of the box height for a chair, 0.25 to 0.65 for a couch and 0.30 to 0.75 for a bed (the templates' own are 0.50, 0.45 and 0.55). The fitted shape is the template stretched in height, piecewise linear, so that the floor, the seat and the top of the box stay where they are and every rectangle of the template stays a rectangle
Other named categories	The fixed template, now facing the right way
`unknown`, `other`, and any template that cannot be built	A shape cut from the scan: the profile reduced to 8×8 (the highest of each 2×2), unseen cells filled from their seen neighbours, one column per cell from the floor to its height, the top and the steps between columns as faces. If that passes 1,500 triangles it is reduced to 4×4. Tinted by category as the box was
No profile, or coverage under 0.30, and no template	The placeholder box. This is the only case left in which a box is seen

The same code builds first shapes on the phone and on the Spectator; the Spectator already reads the whole RoomModel from the Forge (`SpectatorRooms.cs`), so it has the profile.

5.3 Area C · The Forge fits, colours and levels the mesh

The job's `meta` gains `profile` (as in the RoomModel) and `triTargetLight`. Clean-up becomes, in order:

1. Weld, decimate to the full target, drop detached fragments (as today, with the new target).
2. **Orient by the fit.**
3. Scale to the box, at most $\pm 15\,\%$ of stretch per axis (as today).
4. Regrid so no edge exceeds 0.25 m and the count ends at or below the target (as today).
5. **Colour per face.**
6. Bake face shade $\times$ ambient occlusion (as today).
7. **Make the light copy** from the full one: decimate to the light target, then steps 4 to 6.
8. Validate both, publish both.

Clean-up now casts rays and makes two copies, so its time-out rises from 20 s to 60 s.

The fit

- **Candidates.** The up axes the Forge tries today, times four yaws (0°, 90°, 180°, 270°), keeping those whose bounds match the box within the stretch limit.
- **Score.** For a candidate, the mesh's own profile is made with the same downward rays as section 5.2. The score is the mean of $|\text{generated} - \text{scanned}|$ over the cells the scan saw, in fractions of the box height; a cell where the mesh has no surface counts as height 0. Lower is better; 0 is a perfect match.
- **Choice.** The lowest score. When the two best differ by less than 0.02, `triposg_front` breaks the tie: `canonical` prefers yaw 0°, `view` prefers the yaw implied by view 0's `camToBox`, `unknown` keeps the lower score.
- **No profile.** `triposg_front` alone decides among the candidates: yaw 0° for `canonical`, the camera's yaw for `view`. With `unknown`, today's rule.
- **Report.** `fit`, `fitYaw`, `fitUp` and `attempts` are added to the clean-up report and to `JobStatus`.

Attempts (decision M7)

A fit above **0.20** is poor. The Forge then generates again from view 1, then from view 2, when those exist, and publishes the attempt with the lowest score. A mesh is always published (decision M6); a poor fit is reported, never hidden. At most three generations per object.

Colour per face (decision M8)

- For each face and each view: the face is seen when it turns towards the camera (cosine above 0.2), its centre projects inside the crop, the matte is solid there, and a ray from the camera to the face centre meets no other triangle first. The crop's own `intrinsics` and `camToBox` give the projection; both are already in the job.
- The face's colour is the mean of the 5×5 pixels around the projected centre, weighted by the cosine, over the views that see it.
- It is mapped to the nearest object cell of the palette (cells 64 to 127), with today's `nearest_entry`.
- Smoothing: three passes in which a face takes the cell that two of its three neighbours share; then regions under 2 % of the surface merge into their largest neighbour, and the smallest regions merge until at most 6 cells remain.
- A face no view saw takes the cell of the nearest seen face along the surface. With no seen face at all, the object keeps one cell from the dominant colour, as today.
- Vertices are split only where two faces of different cells meet, and every vertex of a face points at the centre of that face's cell. The shader, the palette and the vertex format do not change. `AssetReady.palette` reports the cell that covers the most surface.

Four departures in `forge/colour.py` (a module with tests; not yet wired into clean-up). Speckle is counted after the three passes and the 2 % merge, not after the limit to 6 cells, because counted after the limit a clean object of eight equal stripes limited to three cells looks speckled and collapses to one cell. A face no view saw takes the cell of the nearest seen face by distance along the surface, centre to centre over faces that share a corner or an edge, not by a count of edges crossed, because the count ran the border between two colours at 45° across an unseen back. After the unseen faces are filled, the 2 % merge and the limit run a second time over the whole mesh, so a small island of seen faces surrounded by unseen ones cannot survive. And "no other triangle first" is an exact ray test from the camera to each face centre, with triangles binned by their footprint in the picture, not a depth buffer (tolerance 0.01 m; a view that would need more than 12,000,000 ray-triangle tests sees nothing).

Two levels (decision M10)

- The full copy meets the object's full target (section 5.4); the light copy meets the room's light target.
- `AssetReady` gains `lightSha256`, `lightUrl`, `lightBytes` and `lightTris`. A reader that ignores them shows the full copy everywhere, as today.
- `objects[k].assetHashLight` records the light copy beside `assetHash`.
- The phone downloads the full copy first, since the walker is in Home when meshes arrive, then the light one.

5.4 Area D · The furniture budget

The rule (decision M9)

One rule, in whole numbers, so that the phone (`ObjectMeshRules`), the Forge (`budget.py`) and the reference (`roommath.py`) give the same targets:

```
N      the objects that receive meshes: the 32 largest by volume, ties to the lower id
a_i    2 * (w*h + h*d + w*d), with w, h, d in whole millimetres: mm(v) = floor(v*1000 + 0.5)
A      the sum of a_i
t_i    clamp( floor(120000 * a_i / A), 3000, 15000 )               the full target of object i
if the sum of t_i is above 120000:
    E = sum of (t_i - 3000)        B = 120000 - 3000*N
    t_i = 3000 + floor( (t_i - 3000) * B / E )
light  clamp( 60000 div N, 1500, 5000 ), whole-number division, the same for every object
```

Triangles are shared by surface area, because triangles cover surface. The total never passes 120,000 and no object falls under 3,000. Objects beyond the 32 largest keep their first shape and have `triTarget` 0.

Objects of equal size	Full target each	Full total	Light target each	Light total
5	15,000	75,000	5,000	25,000
8	15,000	120,000	5,000	40,000
12	10,000	120,000	5,000	60,000
20	6,000	120,000	3,000	60,000
32	3,750	120,000	1,875	60,000

What each copy must satisfy

	Full copy	Light copy	Today
Triangles	1,500 to 15,000, and at most the object's target	1,500 to 5,000, and at most the light target	1,500 to 5,000
File	At most 2 MB	At most 1 MB	1 MB
Longest edge	0.25 m	0.25 m	0.25 m
Palette cells	One or more object cells	One or more object cells	One
Everything else	As today: one mesh, no material, no texture, 16-bit indices, origin at the base centre, within 10 % of the box		

The worst full copy, every vertex split, is 45,000 vertices of 36 bytes and 90,000 bytes of indices: 1.71 MB.

Which copy a space shows

Every space states the level it wants; the room content holds both meshes per object and swaps when a space is entered, never during a frame.

Space	Level	Room cost per frame, both eyes	New planning row
S0 Home	Full	246,000 triangles, 68 draw calls	At most 80 draws and 250k triangles
S2 Curved Walk	Full	246,000 triangles	At most 240 draws and 370k triangles
S1, S3, S4, S5, S6	Light	126,000 triangles, as today	Triangles as today; 24 more draw calls each, for the 12 more objects

The must-limits of the plan (300 draw calls and 500,000 triangles per frame) do not change, and every row stays inside them. Of these rows only Home's is held by code (`RoomBudget`: 80 draw calls and 250,000 triangles per frame, counted from what the builders make); the others are planning figures until check M3. Three further controls:

- **The quality ladder.** The quality table gains one column, furniture level. It is a table in `Rendering/Quality/QualityFurniture.cs`, not a field of `QualityRow`, which lives in `Atlas.World`. Nominal and High allow Full; Fair, Serious and Critical force Light in every space, so the first step below nominal drops the furniture to its light copies before any step lowers the render scale. `QualityApplier.FurnitureCap` and `QualityLadder.FurnitureCap` publish the column, and `FurnitureLevelPolicy` (`App/Spaces/FurnitureLevels.cs`) puts the resulting level on the room when a space is entered, when the tier changes, when the parameter changes and when the room is rebuilt.
- **A parameter.** `render.furniture` = `auto`, `light` or `full`, so the Console can force a level: `light` forces Light in every space, `full` forces Full whatever the ladder says.
- **The Spectator** always shows the full copy: it downloads no light copy and sets no level.

The measurement (check M3)

A new fixture, `budget_room_120k`: the room of `budget_room` with 24 stand-ins of exactly 5,000 triangles, 120,000 in all. On the phone, Home and Curved Walk are run through the existing PerfTour with it. A space that does not hold 60 frames per second there shows the light copy by default; the full copies are still made for the spaces that hold. Until the measurement exists, the defaults of the table stand.

5.5 Area E · Doors and windows

The shell stays as it is: the floor from the confirmed outline, the walls on the outline's edges. That is what makes a virtual wall coincide with the real one. The wall bridge (`P1-C05`), already written and tested headless, adds what the shell lacks:

1. **Build the plugin.** A post-build step, `Editor/AtlasRoomPlanProject.cs` (called from `AtlasPostBuild`), adds `Plugins/iOS/AtlasRoomPlan.swift` to the sources of the `UnityFramework` target of the generated Xcode project, links the RoomPlan framework weakly and sets the Swift build settings; applied twice it equals applied once. What was proved: the step was applied to a copy of an earlier Unity export (made on 3 October 2026, without the Swift file) and the `UnityFramework` target then built (`BUILD SUCCEEDED`), exporting the four `AtlasRoomPlan` functions with RoomPlan linked weakly. The app target (`Unity-iPhone`) stopped at its two storyboards, "iOS 26.1 Platform Not Installed": the iOS platform component is not installed on the build Mac. So it is proved that the file and the settings build in a framework target, not that Unity's export of this tree does, and nothing has run on a phone.
2. **Switch it on.** `WallBridge.Apply(bool)` sets `WallBridge.Enabled` to the parameter `scan.wallBridge` (default on; it is the kill switch) and the bridge's own report that it is available. The parameter is not in the on-device menu's kill-switch list, because the link contract fixes that list as the 21 names of its section 1.5 and a test asserts the count; it is on the menu's Parameters page.
3. **Nothing else moves.** A RoomPlan wall still only matches an outline wall; the outline wall keeps its end points. It takes the height, and its doors, windows and openings. With no capture, an unreadable file or no match, the walls are as today.

The room's distance field treats every outline edge as a wall, with or without an opening (rule G7), so doors do not change where the walker may go. `RoomShellBuilder` already cuts openings into walls and kerbs.

5.6 Area F · Proof

The record (decision M14)

One scan of a real room by the team, taken in a session with `FORGE_DEBUG_KEEP=1`. Under the Forge's data folder, which is already outside git (`station/forge/.gitignore`), it holds: the RoomModel with its profiles, the crops, the generator's raw output per attempt, the job reports, the published GLBs, and `truth.json`, written once by hand: for each object, which side is really the front. Because the raw generator output is kept, clean-up can be changed and run again without the GPU.

`tools/meshcheck/check.py` reads the record and prints the numbers below. Only those numbers, and pictures rendered from the meshes, go into `docs/evidence/mesh/`. No camera frame leaves the Station.

The checks

Check	Measured on	Passes when
M1 First shapes	The record, and the phone	No object with a profile is a plain box. Facing is right (the correct quarter turn) for at least 90 % of chairs and couches and for every wall-backed object. On the phone the first shapes are there within 10 s of entering Home
M2 Generated meshes	The record	Each of the 32 largest objects has a full and a light copy that pass the validator. Generation takes 60 s or less per attempt. The median fit is 0.12 or better. Facing is right for at least 90 % of objects that have a front. Every object above 0.20 after its attempts is listed by name
M3 Budget	The phone, `budget_room_120k`	Home and Curved Walk hold 60 frames per second with the full copies, or the space's default is changed to light and the table of section 5.4 is corrected
M4 Doors and windows	The phone	In a room with a door and a window, both appear in `walls[].openings`; touching each side of the door frame is within 5 cm of the drawn edge

The walk-through on the phone

- [ ] Every chair faces the way the real chair faces.
- [ ] Nothing in the room is a plain box.
- [ ] Each first shape is replaced by a generated mesh, largest objects first.
- [ ] The generated chair is recognisably that chair, not a chair.
- [ ] Touch points on a table edge and a chair back are as close as they were before this work.
- [ ] An object with two clearly different colours shows both.
- [ ] Entering Aligned Cells and coming back to Home changes the meshes' detail and nothing else.
- [ ] The door and the window are where the real ones are.

6 Ray casting

6.1 Where rays are used in this design

Use	Rays	Against	Runs on
The object profile	1,024 per object, straight down	The LiDAR mesh inside the box	The phone, at the end of the scan
The fit	The same 1,024, per candidate orientation	The generated mesh	The Forge
Colour, is this face seen	One per face and view, from the camera	The generated mesh	The Forge
Ambient occlusion	32 directions per vertex (exists today)	The generated mesh	The Forge

None of these runs in a frame. They decide facing, orientation and colour; they do not add detail to a mesh.

6.2 A ray-cast renderer: a future option

Not built now. This is what it would be and what would have to be true first.

What exists. Curved space is drawn by moving every vertex with the warp, which places it by the geodesic from the head: its direction and its length. A straight line of sight from the head is therefore already a geodesic. At every vertex the picture is exactly what a ray caster would draw. The only error is along an edge between two vertices, which is drawn straight where the true edge bends, and the 0.25 m edge rule bounds it.

What a ray caster would bring. No edge rule, so no triangles spent only on bending: a table top would no longer need a grid. Exact depth order on the sphere, where the far side wraps round. It would not add detail to any mesh, and it would not change where anything stands.

What it costs. Every pixel follows a geodesic into the scene, for two eyes, 60 times a second, on a phone. That is practical for content given by a formula (the floor and the tiling) and for distance fields. For triangle meshes it needs a search structure in the shader or hardware ray tracing, and whether Unity offers that on the target phones has not been checked.

A path in three steps, each useful alone.

1. **A reference ray caster in the Editor**, on the processor: the exact picture of a test room. It measures the warp's error in pixels against edge length and distance. If the error is under a pixel beyond some distance, the edge rule can be relaxed there, which saves triangles with no renderer at all.
2. **The floor and the tiles per pixel**, in the fragment shader: the tiling is a formula, so a pixel can find its tile without a mesh.
3. **Furniture as distance fields**, made by the Forge as a third copy, traced per pixel.

Before starting. Check M3 has shown spare frame time; step 1 has shown an error a person can see; and the support for step 3 on the phones has been checked.

7 Classification

Nothing changes (decision M12). For the record, this is who classifies what:

What	Classified by	Classes	Used for
Triangles of the scan	ARKit's mesh classification	wall, floor, ceiling, table, seat, window, door, none	Free floor, obstacles, wall directions, and now the object profile (floor triangles are left out)
Triangles when the mesh has no classes	Our height and normal test (`ScanSolver.cs:321`)	floor, wall, ceiling, table, none	The same, on a phone without LiDAR classes
Objects	RoomPlan, through AR Foundation's bounding boxes	23 categories, 16 of which RoomPlan supplies	The tint, the first shape, the facing rule
Obstacle blobs with no RoomPlan box	Nobody: `unknown`	–	A box today; a scan-cut shape after this design

TripoSG needs no category: it works from the photo. So an `unknown` object gets a generated mesh like any other. Facing is geometry, not classification.

8 What changes in the frozen contracts

The contracts were frozen at G0. This design changes them; every edit is listed.

Document or file	Section	Change
`documents/ATLAS_Master_Plan.md`	D20, 4.5	32 objects, 120,000 triangles, the rule of section 5.4, first shapes, two levels
	6.4 asset contract	The table of section 5.4
	7.4	Clean-up order of section 5.3; the fit; attempts; colour per face; `AssetReady` light fields
	8.5	The rows of section 5.4; the ladder's furniture column
`docs/contracts/lane-c-room.md`	2.1, 2.4	Optional `profile`, `facing`, `assetHashLight`; `triTarget` 0 or 3,000 to 15,000; local +Z is the front
	2.7	R17 follows the new rule; new R21: a profile is exactly 256 bytes
	2.8	New rule G14: the profile and facing, after G13
	9.1, 9.4	Two meshes per object and a level; first shapes
	10.1	Route 4 is on when available
	11.2	A view from the back half of the object, or from more than 30° above it, scores half
`docs/contracts/roommodel.schema.json`	`objects`	The optional keys; the `triTarget` range
`docs/contracts/lane-d-link.md`	2.5	`AssetReady` light fields
	3.6 to 3.9	`meta.profile`, `meta.triTargetLight`; matting required and the picture framed; the fallback decoder; the fit; attempts; the report; cleaning time-out 60 s
	4.1, 4.2	The table of section 5.4 in both validators
	4.3, 4.4	First shapes shown at once; fitted and scan-cut shapes; two downloads per object
`docs/contracts/lane-a-platform.md`	10	An asset may use several object cells
	Quality table	The furniture column; `render.furniture`

The constants that carry the old numbers: `station/forge/forge/budget.py`, `station/forge/forge/validator.py`, `World/Room/RoomModelValidator.cs:15`, `Telemetry/Assets/GlbValidator.cs:37`, `Rendering/Room/RoomBudget.cs:32`, `tools/rooms/roommath.py`, `tools/rooms/validate_room.py`. The five fixtures are regenerated by `tools/rooms/make_fixtures.py`; their geometry hashes do not change, because `triTarget`, `profile` and `facing` are not in the hash.

9 Risks

Risk	What happens then
TripoSG cannot be made to run on the Station, or is too slow	The number of steps is a setting (the official demo allows 8 to 50). If it still does not hold, another generator goes behind the same interface: Hunyuan3D-2 states 6 GB for shape and runs on Windows; TRELLIS (16 GB) and SAM 3D Objects (32 GB, Linux) do not fit the card. Until then the first shapes of area B are what the walker sees, which is already better than today
`diso` does not build on the Station, or fails as in the open report	The adapter's fallback decoder. If that fails too, the row above
The generated mesh does not look like the object	It is shown anyway (M6). The fit and the attempts pick the best of three. The record lists the poor ones so the cause can be found
The scan of an object is thin (it was only seen from one side)	Coverage is in the profile. Under 0.30 the profile is not used for facing or for the first shape; the next rule in the list answers
RoomPlan's box axis turns out to be reliable for some category	The record shows it per category; that category's list is set to `box`
The photo and the mesh do not line up, so colours smear	Palette steps and smoothing absorb small errors. If an object ends up speckled, it falls back to one cell when more than half its faces changed cell during smoothing
The full copies are too heavy on the phone	Check M3 switches the space to the light copy, which is today's load
The wall bridge cannot share the AR session with the box provider, or does not build	The parameter switches it off and the walls are as today
A second RoomModel reader rejects the new keys	The keys are optional and every reader in the tree ignores unknown keys except `validate_room.py`, which is changed with the schema

10 Open points

1. TripoSG's output frame is inferred, not documented (section 5.1), and is measured on area A's first run. The design works with any of the three settings; the setting only breaks ties.
2. The thresholds are starting values: 0.15 and 0.10 m for facing, 0.20 for a poor fit, 0.30 for coverage, 6 cells per object. Checks M1 and M2 on the record are where they are tuned.
3. Unity's support for ray tracing on the target phones is unchecked. It matters only for section 6.2, step 3.
4. The two profile implementations differ in one way far below anything measurable. The phone's limits are the float constants 0.02f and 0.05f widened to double (0.0199999996 and 0.0500000007), and `forge/profile.py` uses 0.02 and 0.05. A hit within about a nanometre of the 0.02 m limit, or of the ray's start 0.05 m above the top of the box, can count on one side only. The shared test vectors do not reach it.
5. Three Playground tests fail in the Unity Editor on `main`, for a reason outside this work: `Blades/Playground/Shell/Cards/HudPanel.cs:117` calls `Destroy` in edit mode. They are not touched by the mesh design and are listed so that nobody takes them for a failure of it.

Appendix A Tasks for agents

One row per task. "Needs" lists tasks that must be done first; tasks with none can start at once.

Task	Lane	What	Done when	Needs
MQ-D01	D	Install TripoSG from our own pinned list and run one object through the Forge; correct the README's install line; record versions. `station/forge/requirements-triposg.txt` (new), `station/forge/README.md`, `docs/versions.md`	A GLB from a real photo passes the validator. Built so far: the pinned list, the run-book and `tools/first_run.py`; the validator has accepted procedural output only, and `docs/versions.md` is not yet written	–
MQ-D02	D	Adapter: the fallback decoder, the checks for `model_index.json` and `diso`, the model loaded at start, and whatever else the first run shows. `forge/generators/triposg.py`, `forge/service.py`	`GET /health` reports ready; three different photos give three meshes. Built so far: the adapter, tested with a fake pipeline	MQ-D01
MQ-D03	D	Input picture: matting required, cut to the matte, white, padded square. `forge/imaging.py`, `forge/service.py`	A test photo's prepared picture is square with the object centred	–

Task	Lane	What	Done when	Needs
MQ-D04	D	Measure and record the output frame. `forge/config.py`	`triposg_up` and `triposg_front` are set from a non-symmetric test object. Built so far: the two settings (environment variables, defaults `Y` and `unknown`) and the tool that reads the frame	MQ-D02
MQ-C01	C	The object profile in the scan solver. `Scan/Build/ObjectProfile.cs` (new), `Scan/Routes/ScanSolver.cs`	Synthetic rooms give the expected bytes. Built instead of "`roommath.py` agrees": the Forge's `forge/profile.py` and the phone's `ObjectProfile` pass the same three test vectors (A, B and C of the shared brief, S3); `roommath.py` has no profile	–
MQ-C02	C	`profile`, `facing`, `assetHashLight` in the RoomModel, the reader, the writer, the schema, rule R21. `World/Room/RoomModel.cs`, `RoomModelJson.cs`, `RoomModelValidator.cs`, `roommodel.schema.json`, `tools/rooms/`	Round trip on a room with profiles; fixtures still valid	–
MQ-C03	C	The facing rule and the quarter-turn normalisation. `Scan/Build/ObjectFacing.cs` (new), `RoomModelBuilder.cs`	Synthetic chair, stool, storage, deep-appliance and table cases face as designed, with the tie rule of section 5.2	MQ-C01, MQ-C02
MQ-C05	C	Photo selection prefers views from the front half and from no more than 30° above. `Scan/Photo/PhotoScore.cs`	Of a front and a back view of equal quality, the front one is view 0	MQ-C03
MQ-D05	D	Fitted templates for chair, couch and bed; the scan-cut shape. `Telemetry/Assets/FallbackMeshBuilder.cs`, `FittedShapes.cs` (new), `tools/gen_fallback_set.py`	Every category and `unknown` build under 1,500 triangles, no edge over 0.25 m	MQ-C02
MQ-D06	D	First shapes shown at once, on the phone and the Spectator. `Telemetry/Assets/AssetClientCore.cs`, `Spectator/SpectatorFurniture.cs`	With the Forge reachable, no object that has a template or a usable profile is a box after the room is built	MQ-D05
MQ-A01	A, C, D	The triangle rule in all three places, and the validators. `RoomModelValidator.cs`, `forge/budget.py`, `tools/rooms/roommath.py`, `GlbValidator.cs`, `forge/validator.py`	The table of section 5.4 from all three; the fixtures regenerate	–
MQ-D07	D	The fit: the mesh's profile, candidates, choice, report; cleaning time-out 60 s. `forge/cleanup.py`, `forge/profile.py` (new), `forge/models.py`, `forge/config.py`	A test mesh turned by each quarter turn is turned back. Built so far: `forge/profile.py` with its tests (the mesh's profile, the fit score, the turning); the fit is not yet in clean-up	MQ-C02
MQ-D08	D	Attempts on a poor fit. `forge/service.py`	A forced poor fit runs views 1 and 2 and publishes the best	MQ-D07
MQ-D09	D	Colour per face. `forge/colour.py` (new), `forge/cleanup.py`, `forge/palette.py`	A two-colour test object gives two cells in the right places. Built so far: `forge/colour.py` with its tests, which include a two-colour object; it is not yet in clean-up	MQ-D07
MQ-D10	D	The light copy; `AssetReady` light fields; two downloads. `forge/cleanup.py`, `forge/service.py`, `Link/Messages/JsonMessages.cs`, `AssetClientCore.cs`	Both copies validate; an old reader still shows the full one. Built so far: the phone's side (the message fields, the Hub pass-through, the validator levels, the light download) and the Forge's validator and `JobStatus` fields; the Forge does not yet make a light copy	MQ-A01
MQ-A02	A, C	Two meshes per object and a level in the room content; each space's level; the ladder column; `render.furniture`. `World/Room/IRoomContent.cs`, `Rendering/Room/RoomContent.cs`, `Rendering/Quality/`, the space modules	Entering a light space swaps meshes; the ladder's first step forces light	MQ-D10
MQ-A03	A	`budget_room_120k` and the budget tests. `tools/rooms/make_fixtures.py`, `Rendering/Room/RoomBudget.cs`, tests	`RoomBudget.Estimate` gives about 246,000 triangles per frame for Home (tested within 5,000 of it, the shell being a few thousand) and 52 draw calls for the 24 stand-ins	MQ-A01

Task	Lane	What	Done when	Needs
MQ-C04	C	Build the Swift plugin into the Xcode project; switch the bridge on. `Plugins/iOS/AtlasRoomPlan.swift`, an Editor post-build step, `Scan/Routes/WallBridge.cs`, `App/Params/AppParams.cs` (`Atlas.Scan` cannot reach the registry, so the parameter is declared there)	The app builds; `scan.wallBridge` turns it off. Built so far: the post-build step, tested headless, and the `UnityFramework` target built from a copy of an earlier export; the app target stopped at storyboards on this Mac (section 5.5); the parameter is tested headless	–
MQ-F01	D	The record and `tools/meshcheck/check.py`; raw generator output kept under debug-keep. `forge/service.py`, `tools/meshcheck/` (new)	The script prints the numbers of M1 and M2 from a record	MQ-D07
MQ-F02	Lead	Take the record; write `truth.json`; run M1 to M4; fill `docs/evidence/mesh/`	The four checks have a result each	All

Appendix B Formulas

Yaw and the quarter turns. Yaw ψ turns local +Z towards local +X (lane C contract, section 1). If the front is found on local +X, the new yaw is $\psi + \pi/2$; on local $-$Z, $\psi + \pi$; on local $-$X, $\psi - \pi/2$; each wrapped to $(-\pi, \pi]$. On +X and $-$X, `size.w` and `size.d` swap. The profile is turned by the same number of quarter turns so that its row 15 is the front.

Side strips of the profile. Back strip of the candidate front +Z: rows $j = 0$ to 2. Of +X: columns $i = 0$ to 2. Of $-$Z: rows 13 to 15. Of $-$X: columns 13 to 15. A strip's mean is over its seen cells, as height ÷ box height.

The fit score. With S the cells the scan saw, p the scanned profile and g the generated mesh's, both as fractions of the box height, and $g = 0$ where the mesh has no surface:

$$\text{score} = \frac{1}{|S|} \sum_{c \in S} |g_c - p_c|.$$

The view's yaw. With `camToBox` the camera-to-box transform of view 0, the camera stands at c, its translation. The azimuth of the view is $\text{atan2}(c_x, c_z)$ in the box's frame; the nearest quarter turn to it is the candidate the tie-break prefers when `triposg_front` is `view`. Its elevation is the angle of $c - (0, h/2, 0)$ above the horizontal.

Appendix C Where things are

Subject	File
The scan routes, the solver	`ATLAS/Assets/Atlas/Scan/Routes/ScanSolver.cs`, `ScanSession.cs`
RoomPlan boxes to object boxes	`ATLAS/Assets/Atlas/Scan/Routes/ObjectBoxMerge.cs`
The RoomModel	`ATLAS/Assets/Atlas/World/Room/`, `docs/contracts/lane-c-room.md`
Photo capture and what is sent	`ATLAS/Assets/Atlas/Scan/Photo/`, `World/Room/IObjectPhotoSink.cs`, `Telemetry/Assets/ForgeTransport.cs`
The Forge	`station/forge/` (`service.py`, `cleanup.py`, `budget.py`, `validator.py`, `generators/`)
The asset client and the built-in set	`ATLAS/Assets/Atlas/Telemetry/Assets/`
The room as drawn	`ATLAS/Assets/Atlas/Rendering/Room/`
The warp and the room in curved space	`Rendering/Shaders/Resources/AtlasWarp.hlsl`, `Rendering/Tiles/CurvedRoomView.cs`
The wall bridge	`ATLAS/Assets/Atlas/Scan/Routes/WallBridge.cs`, `ATLAS/Assets/Plugins/iOS/AtlasRoomPlan.swift`

Outside the tree (read on 4 October 2026; nothing was run):

Subject	Source
TripoSG: the inference script, the framing, the pipeline, the requirements	`https://github.com/VAST-AI-Research/TripoSG` (`scripts/inference_triposg.py`, `scripts/image_process.py`, `triposg/pipelines/pipeline_triposg.py`, `requirements.txt`)
TripoSG: the weights, their size and licence	`https://huggingface.co/VAST-AI/TripoSG`

Subject	Source
TripoSG: training views, the 90 % foreground	`https://arxiv.org/abs/2502.06608` (sections 5.1 and 5.2.1)
The decoder failing on Windows	`https://github.com/VAST-AI-Research/TripoSG/issues/63`
The `transformers` version that fails	`https://github.com/VAST-AI-Research/TripoSG/issues/61`
`diso`: built from source, its licence	`https://github.com/SarahWeiii/diso`
Other generators, for the risk table	`https://github.com/Tencent-Hunyuan/Hunyuan3D-2`, `https://github.com/microsoft/TRELLIS`, `https://github.com/facebookresearch/sam-3d-objects`